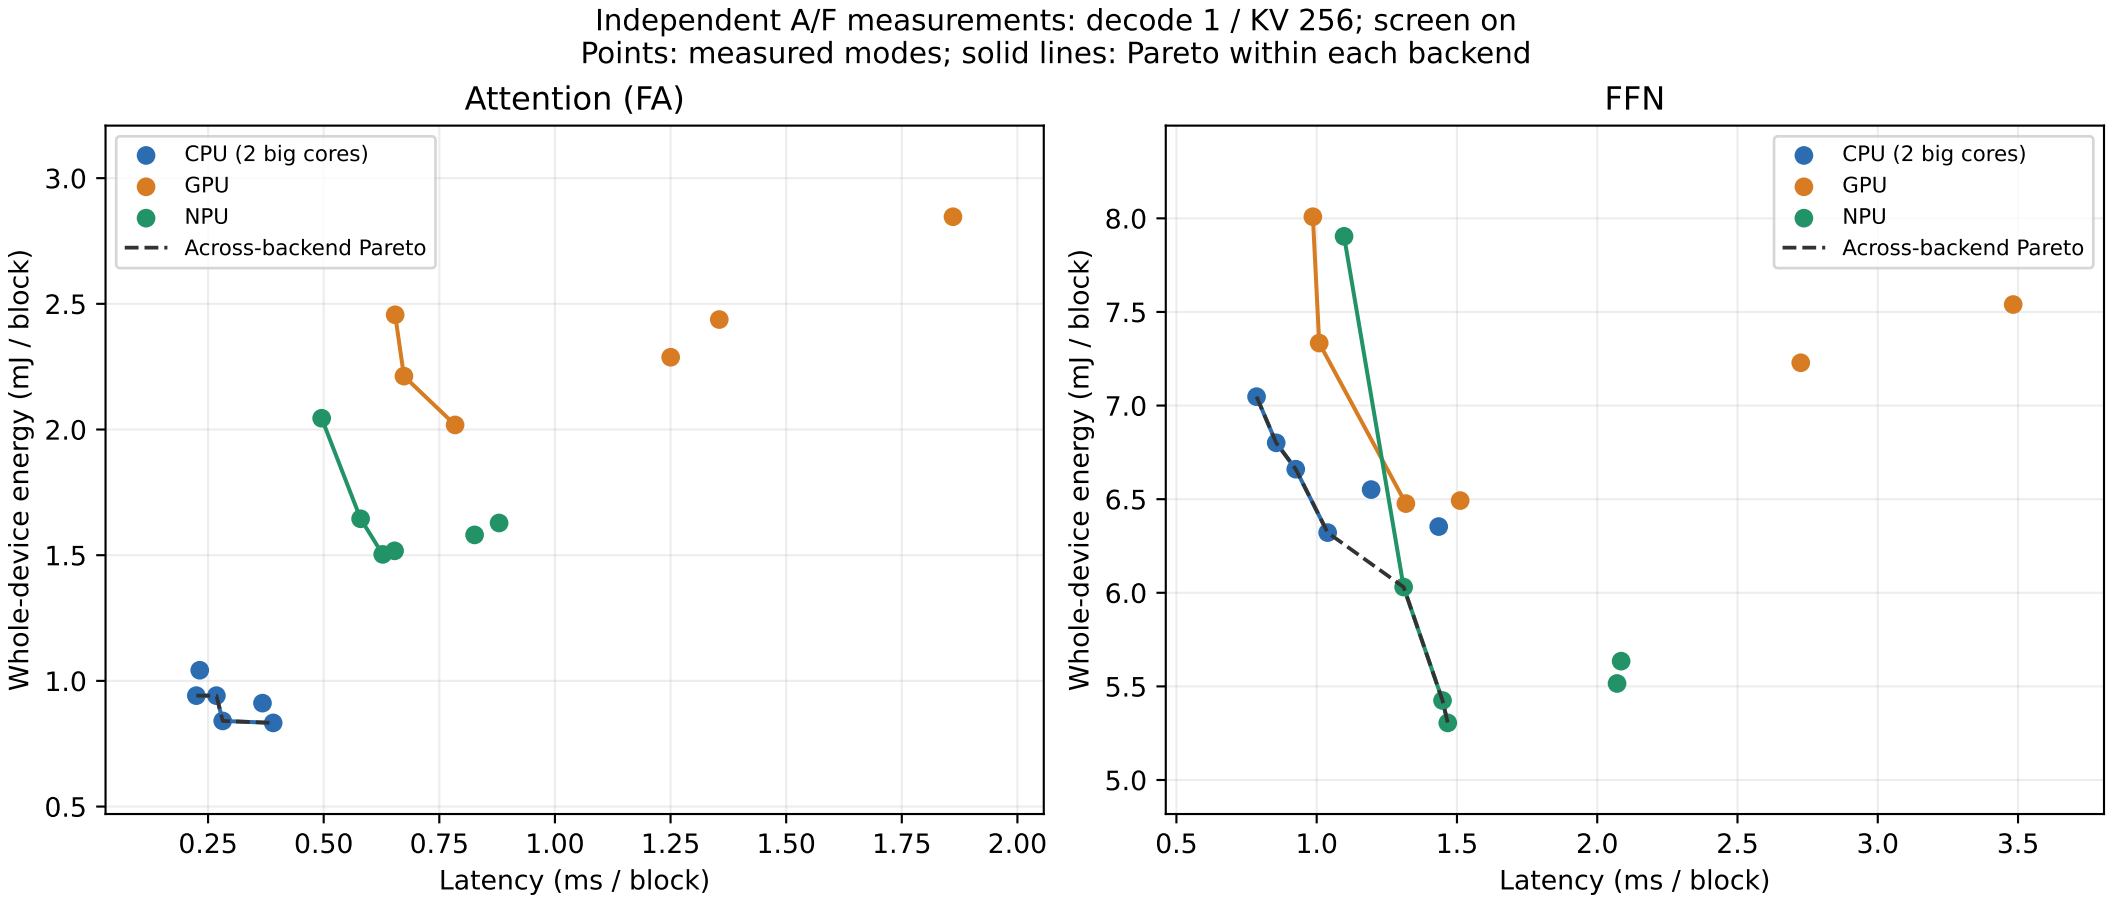
Independent A/F measurements: decode 1 / KV 256; screen on
Points: measured modes; solid lines: Pareto within each backend
Attention (FA)
FFN
CPU (2 big cores)
GPU
NPU
Across-backend Pareto
Whole-device energy (mJ / block)
Latency (ms / block)
0.5
1.0
1.5
2.0
2.5
3.0
0.25
0.50
0.75
1.00
1.25
1.50
1.75
2.00
5.0
5.5
6.0
6.5
7.0
7.5
8.0
3.5