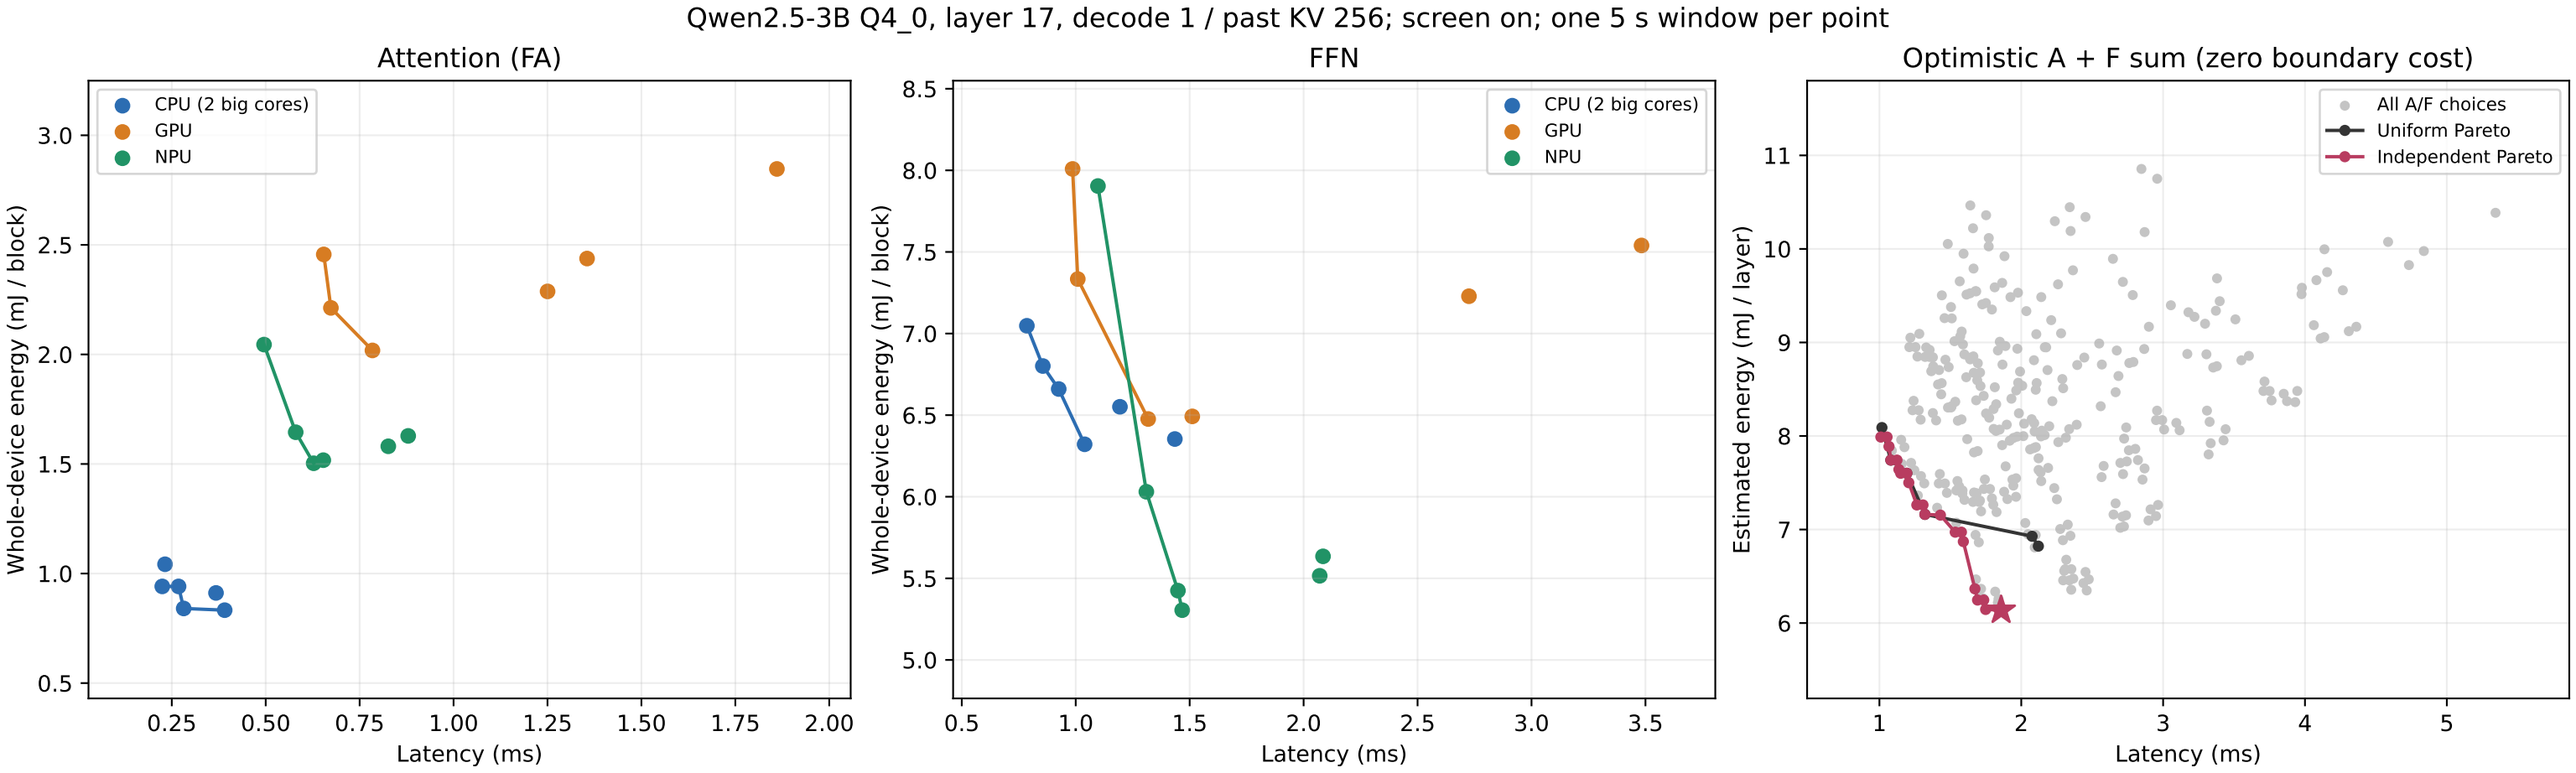
Qwen2.5-3B Q4_0, layer 17, decode 1 / past KV 256; screen on; one 5 s window per point
Attention (FA)
FFN
Optimistic A + F sum (zero boundary cost)
CPU (2 big cores)
GPU
NPU
All A/F choices
Uniform Pareto
Independent Pareto
Whole-device energy (mJ / block)
Estimated energy (mJ / layer)
Latency (ms)