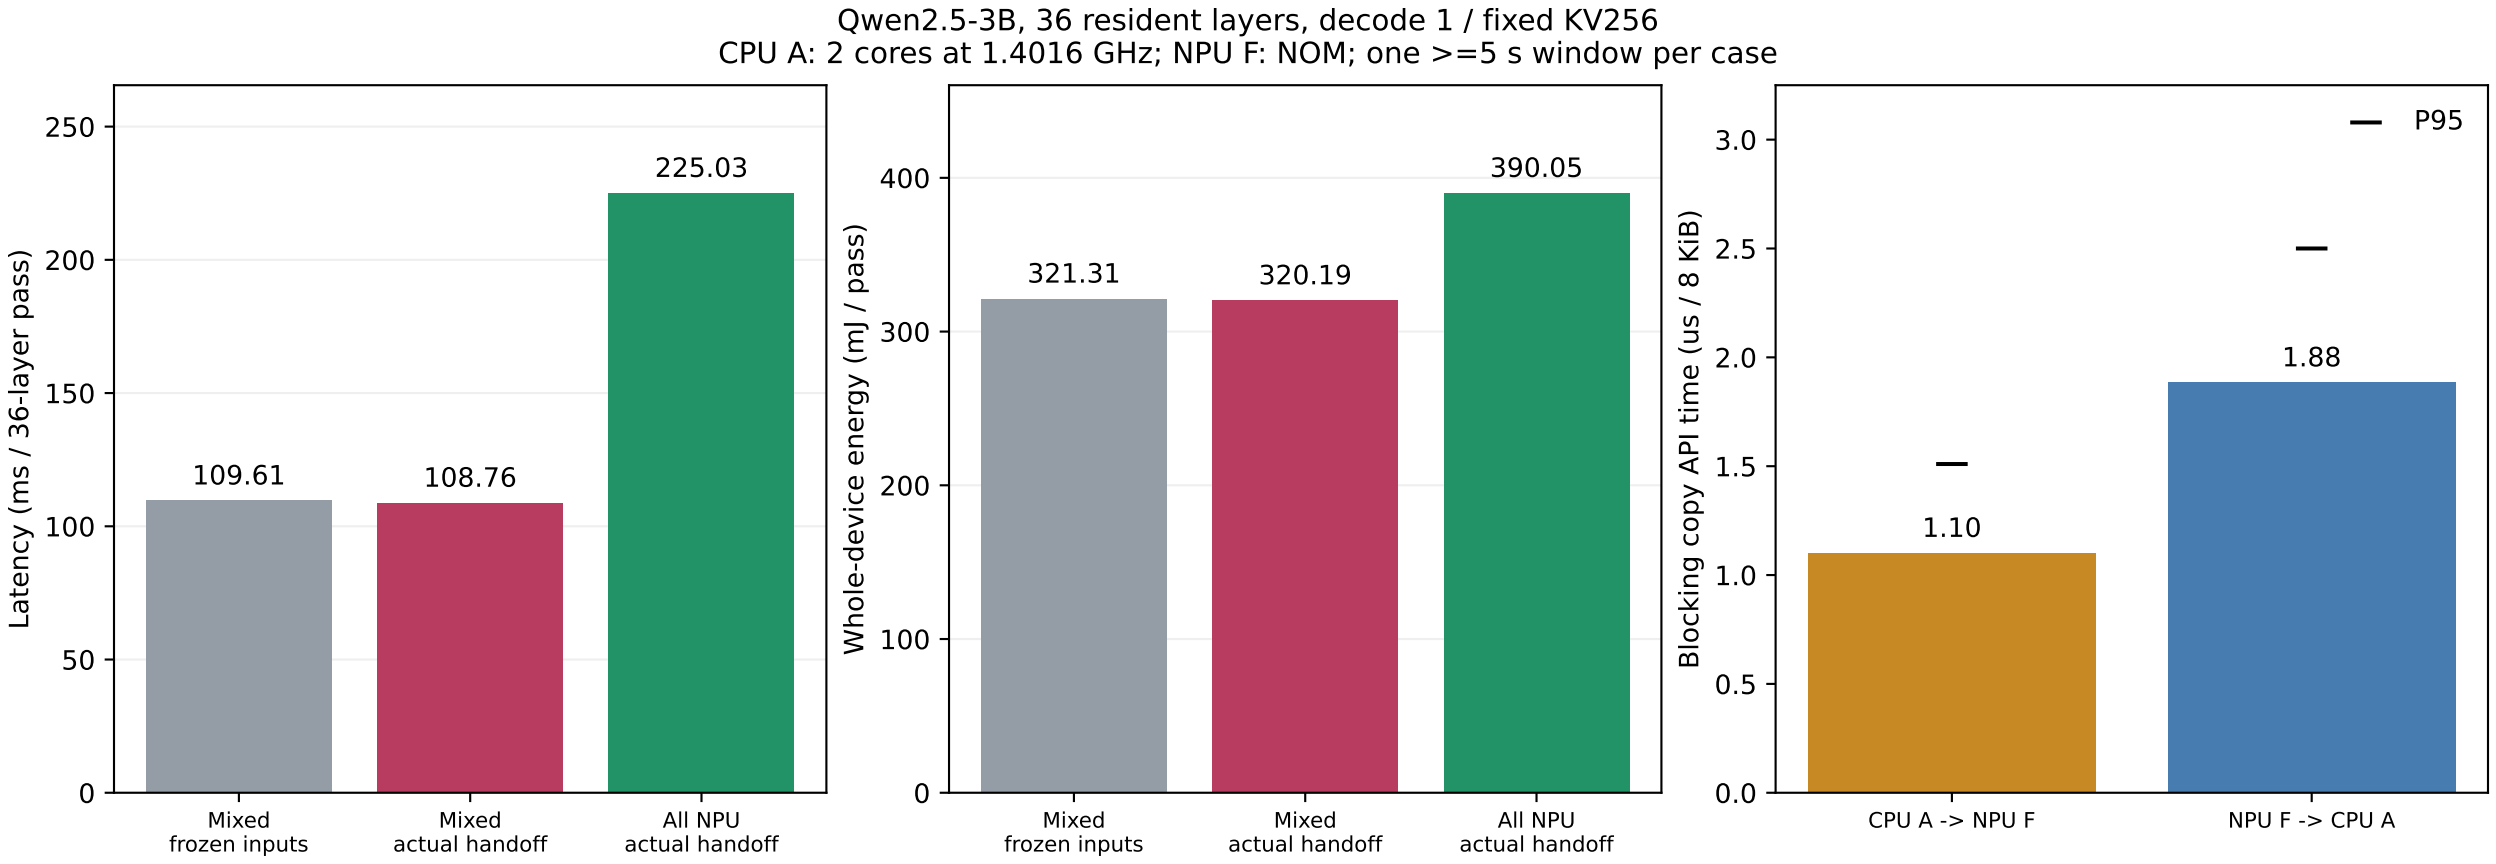
Qwen2.5-3B, 36 resident layers, decode 1 / fixed KV256
CPU A: 2 cores at 1.4016 GHz; NPU F: NOM; one >=5 s window per case
Latency (ms / 36-layer pass)
0
50
100
150
200
250
109.61
108.76
225.03
Mixed frozen inputs
Mixed actual handoff
All NPU actual handoff
Whole-device energy (mJ / pass)
0
100
200
300
400
321.31
320.19
390.05
Mixed frozen inputs
Mixed actual handoff
All NPU actual handoff
Blocking copy API time (us / 8 KiB)
0.0
0.5
1.0
1.5
2.0
2.5
3.0
P95
1.10
1.88
CPU A -> NPU F
NPU F -> CPU A