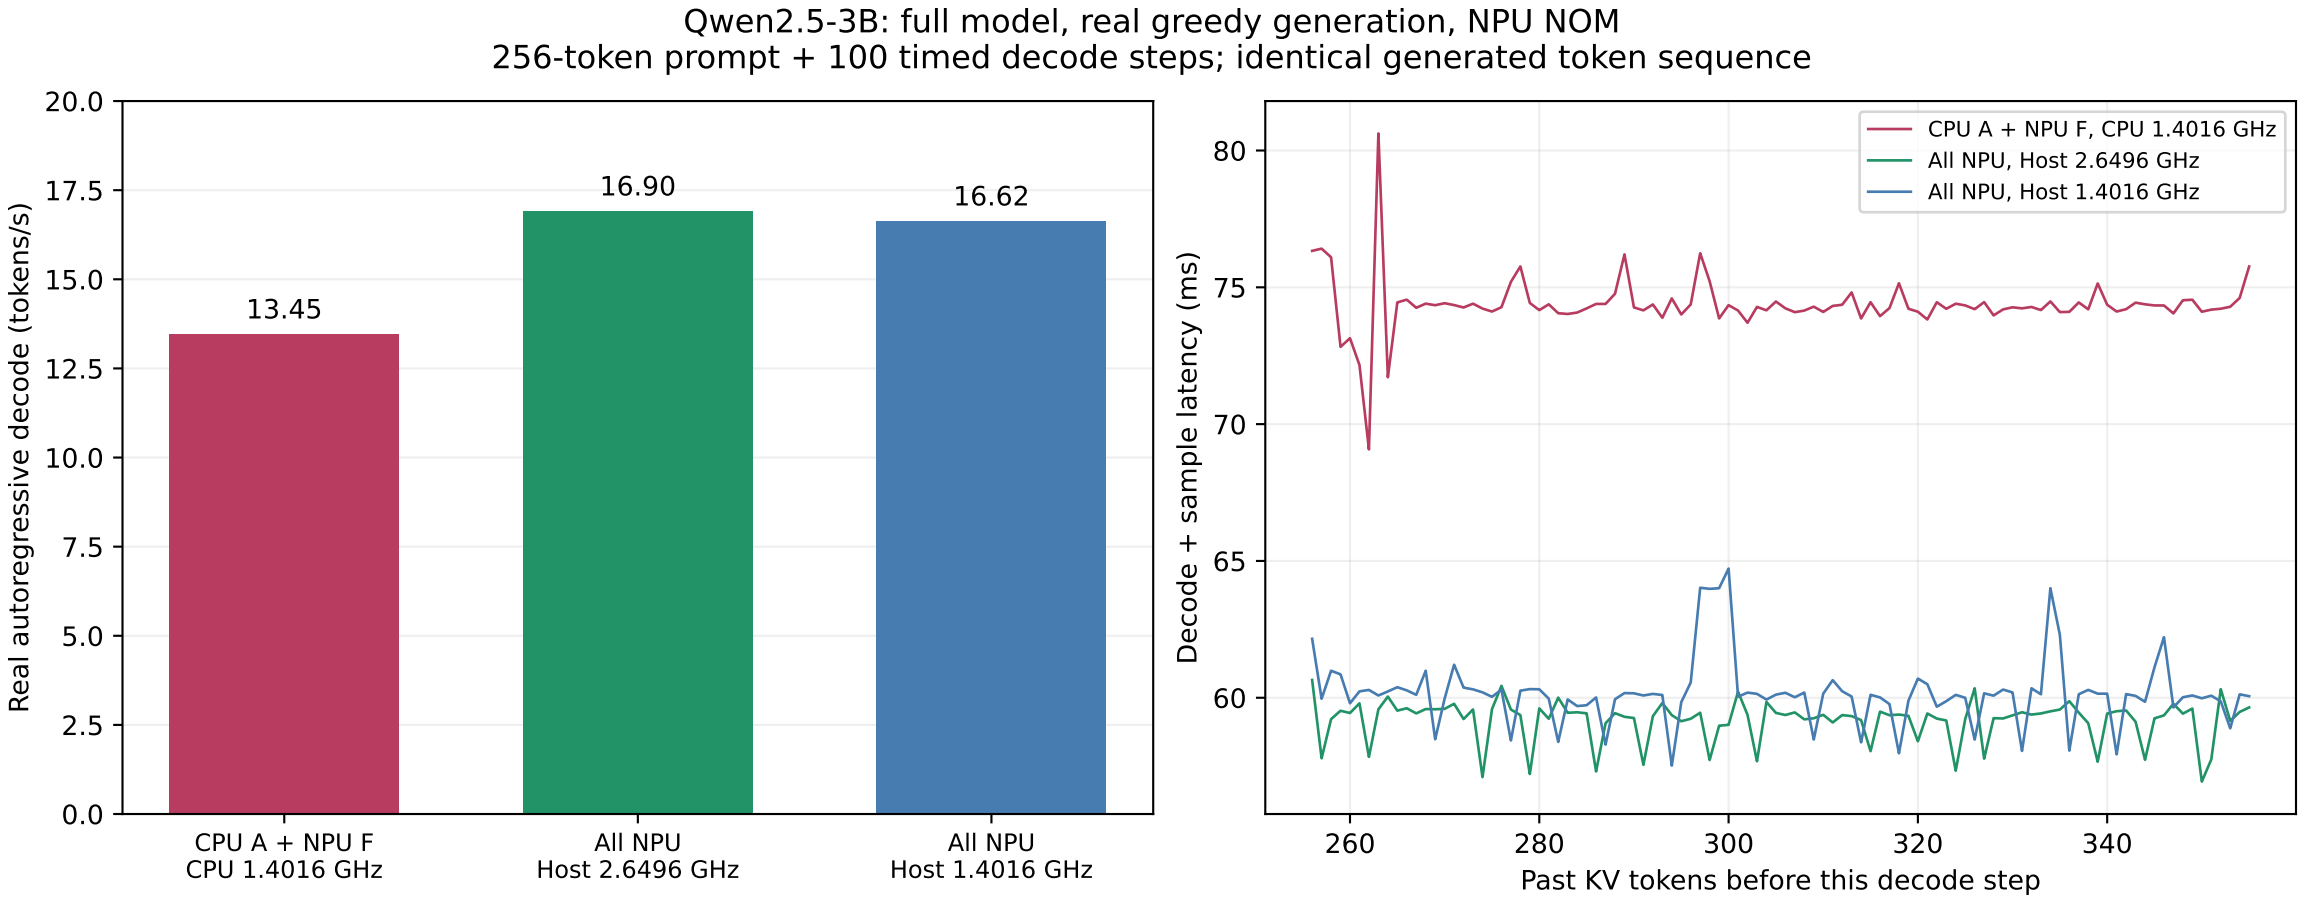
Qwen2.5-3B: full model, real greedy generation, NPU NOM
256-token prompt + 100 timed decode steps; identical generated token sequence
Real autoregressive decode (tokens/s)
20.0
17.5
15.0
12.5
10.0
7.5
5.0
2.5
0.0
13.45
16.90
16.62
CPU A + NPU F
CPU 1.4016 GHz
All NPU
Host 2.6496 GHz
All NPU
Host 1.4016 GHz
Decode + sample latency (ms)
80
75
70
65
60
260
280
300
320
340
Past KV tokens before this decode step
CPU A + NPU F, CPU 1.4016 GHz
All NPU, Host 2.6496 GHz
All NPU, Host 1.4016 GHz