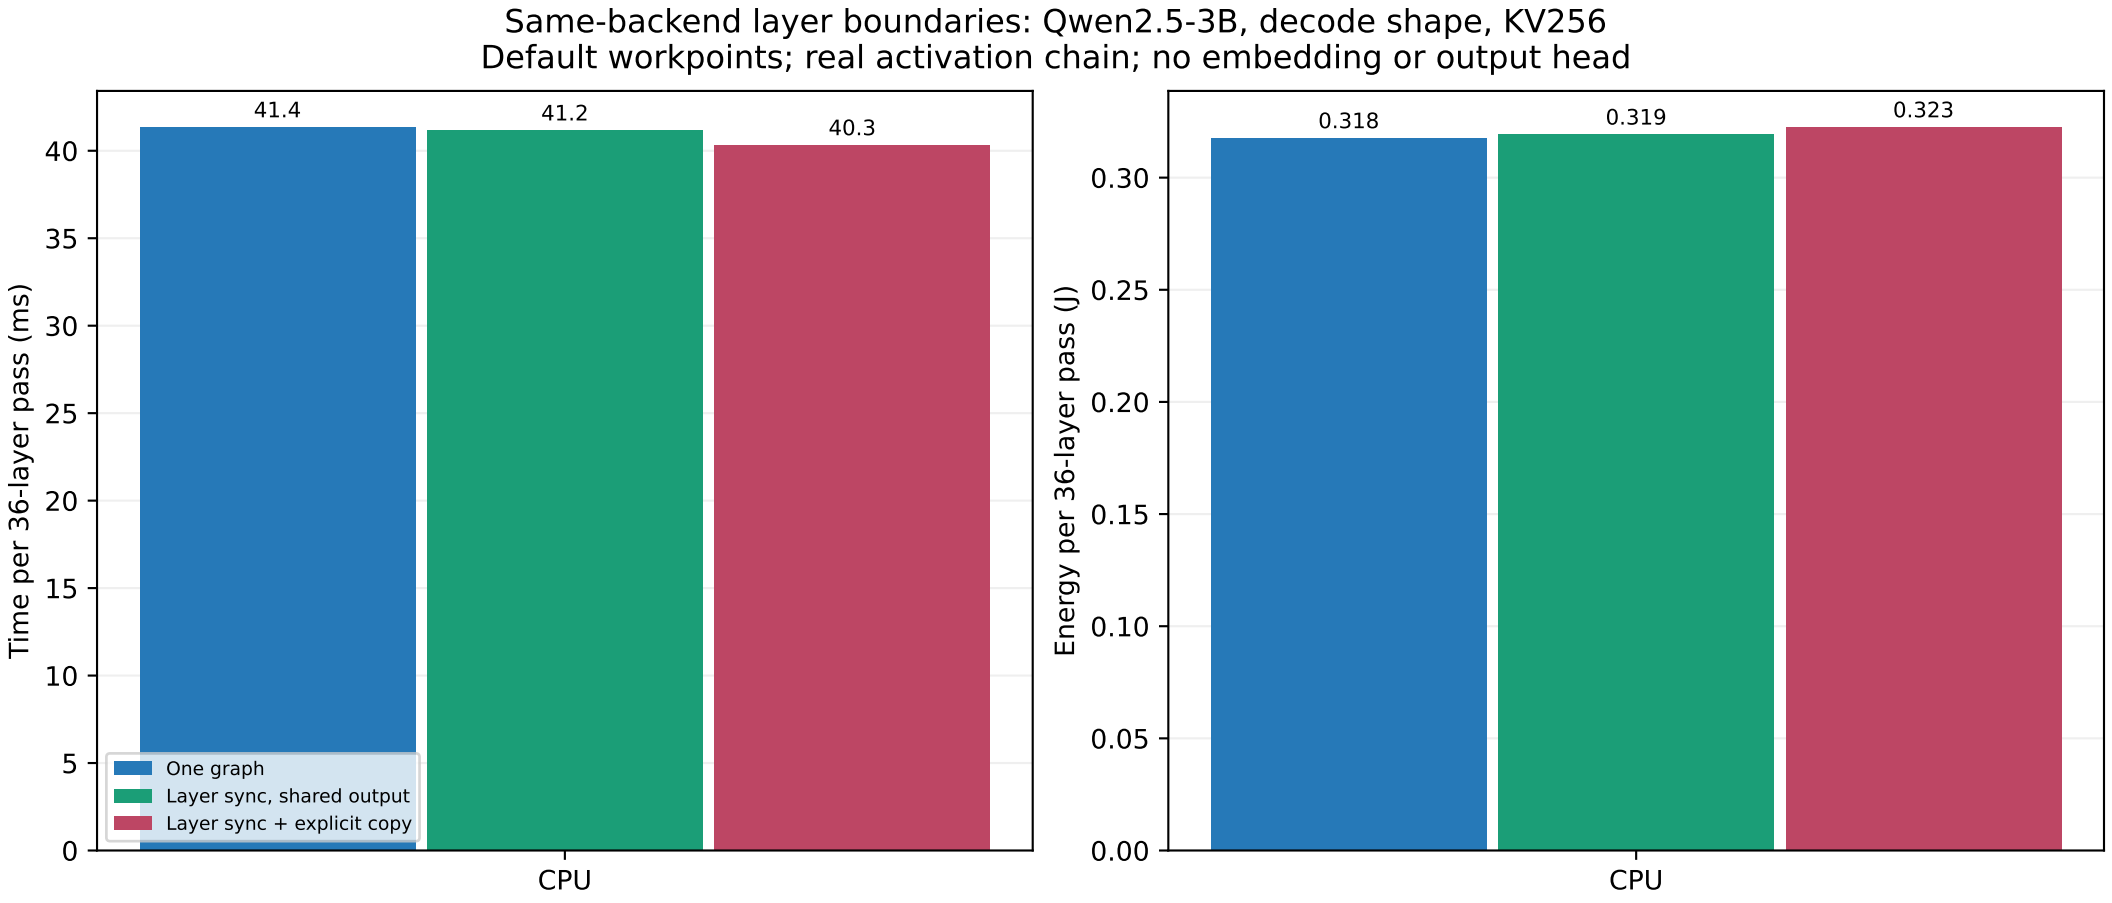
Same-backend layer boundaries: Qwen2.5-3B, decode shape, KV256
Default workpoints; real activation chain; no embedding or output head
Time per 36-layer pass (ms)
41.4
41.2
40.3
One graph
Layer sync, shared output
Layer sync + explicit copy
CPU
Energy per 36-layer pass (J)
0.318
0.319
0.323
CPU