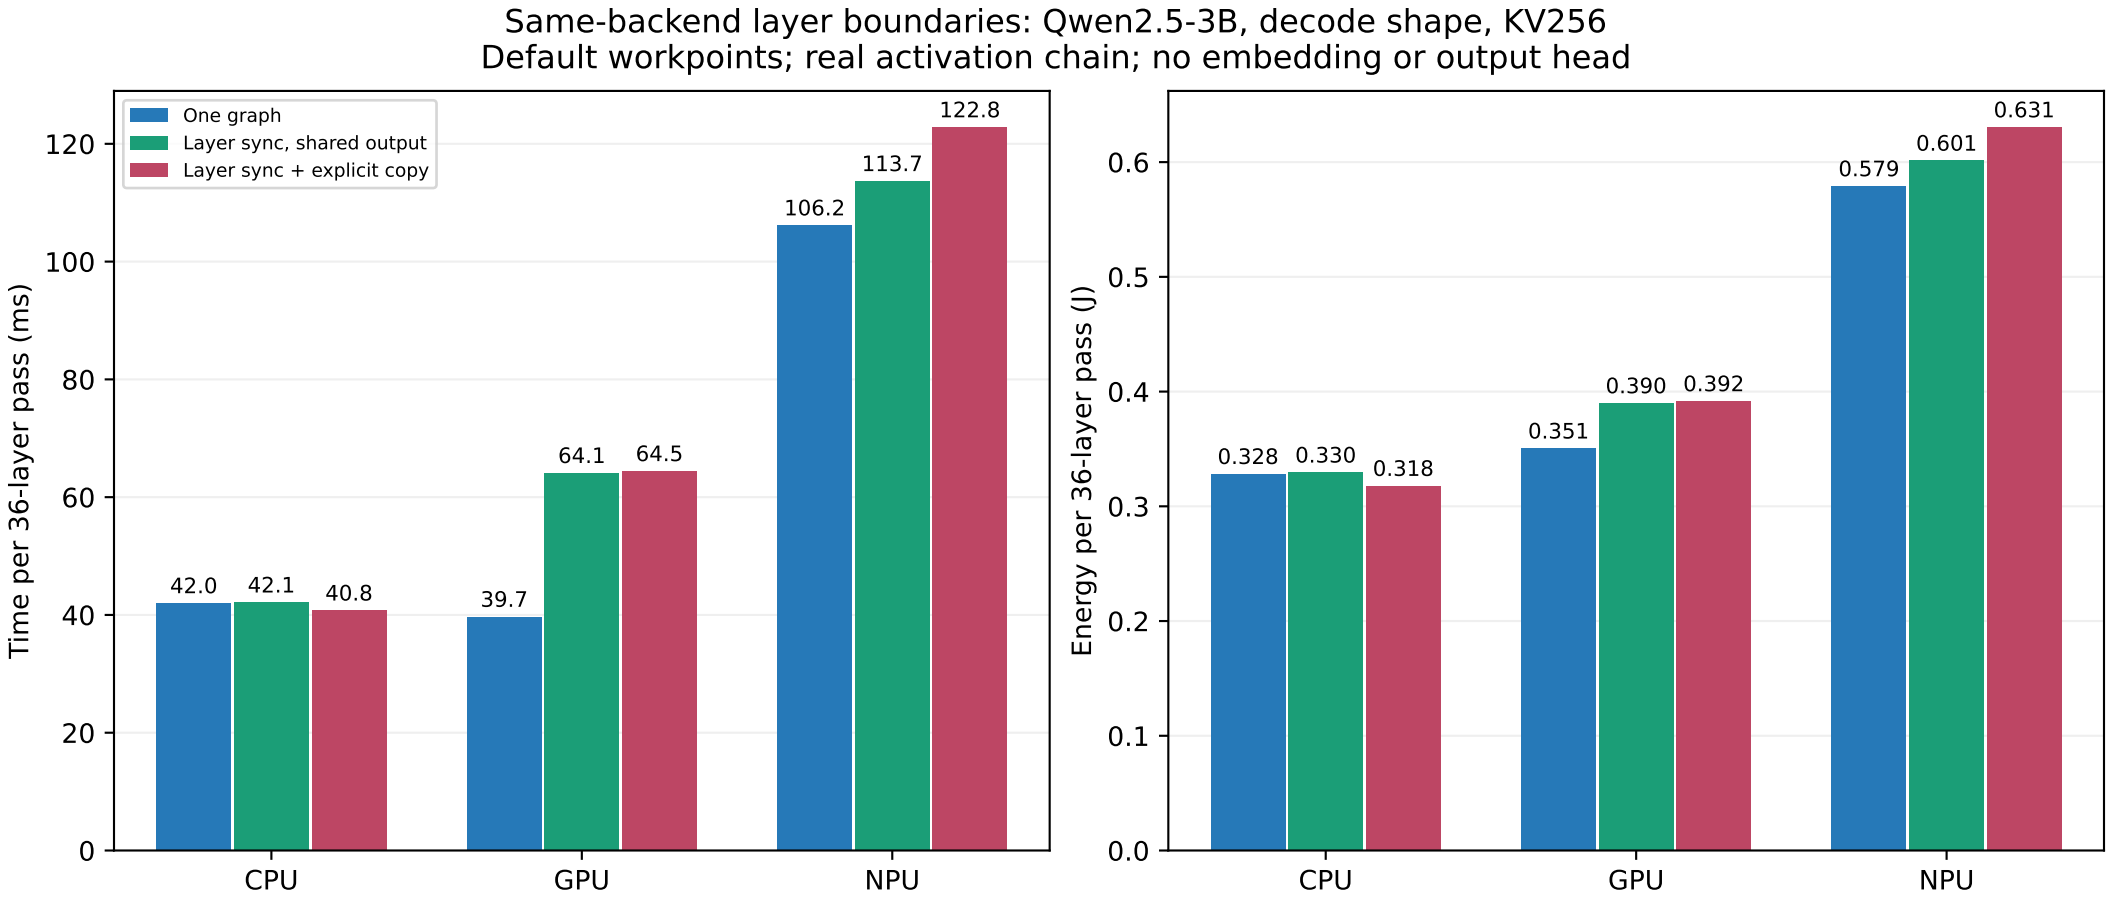
Same-backend layer boundaries: Qwen2.5-3B, decode shape, KV256
Default workpoints; real activation chain; no embedding or output head
One graph
Layer sync, shared output
Layer sync + explicit copy
Time per 36-layer pass (ms)
0
20
40
60
80
100
120
42.0
42.1
40.8
39.7
64.1
64.5
106.2
113.7
122.8
CPU
GPU
NPU
Energy per 36-layer pass (J)
0.0
0.1
0.2
0.3
0.4
0.5
0.6
0.328
0.330
0.318
0.351
0.390
0.392
0.579
0.601
0.631
CPU
GPU
NPU