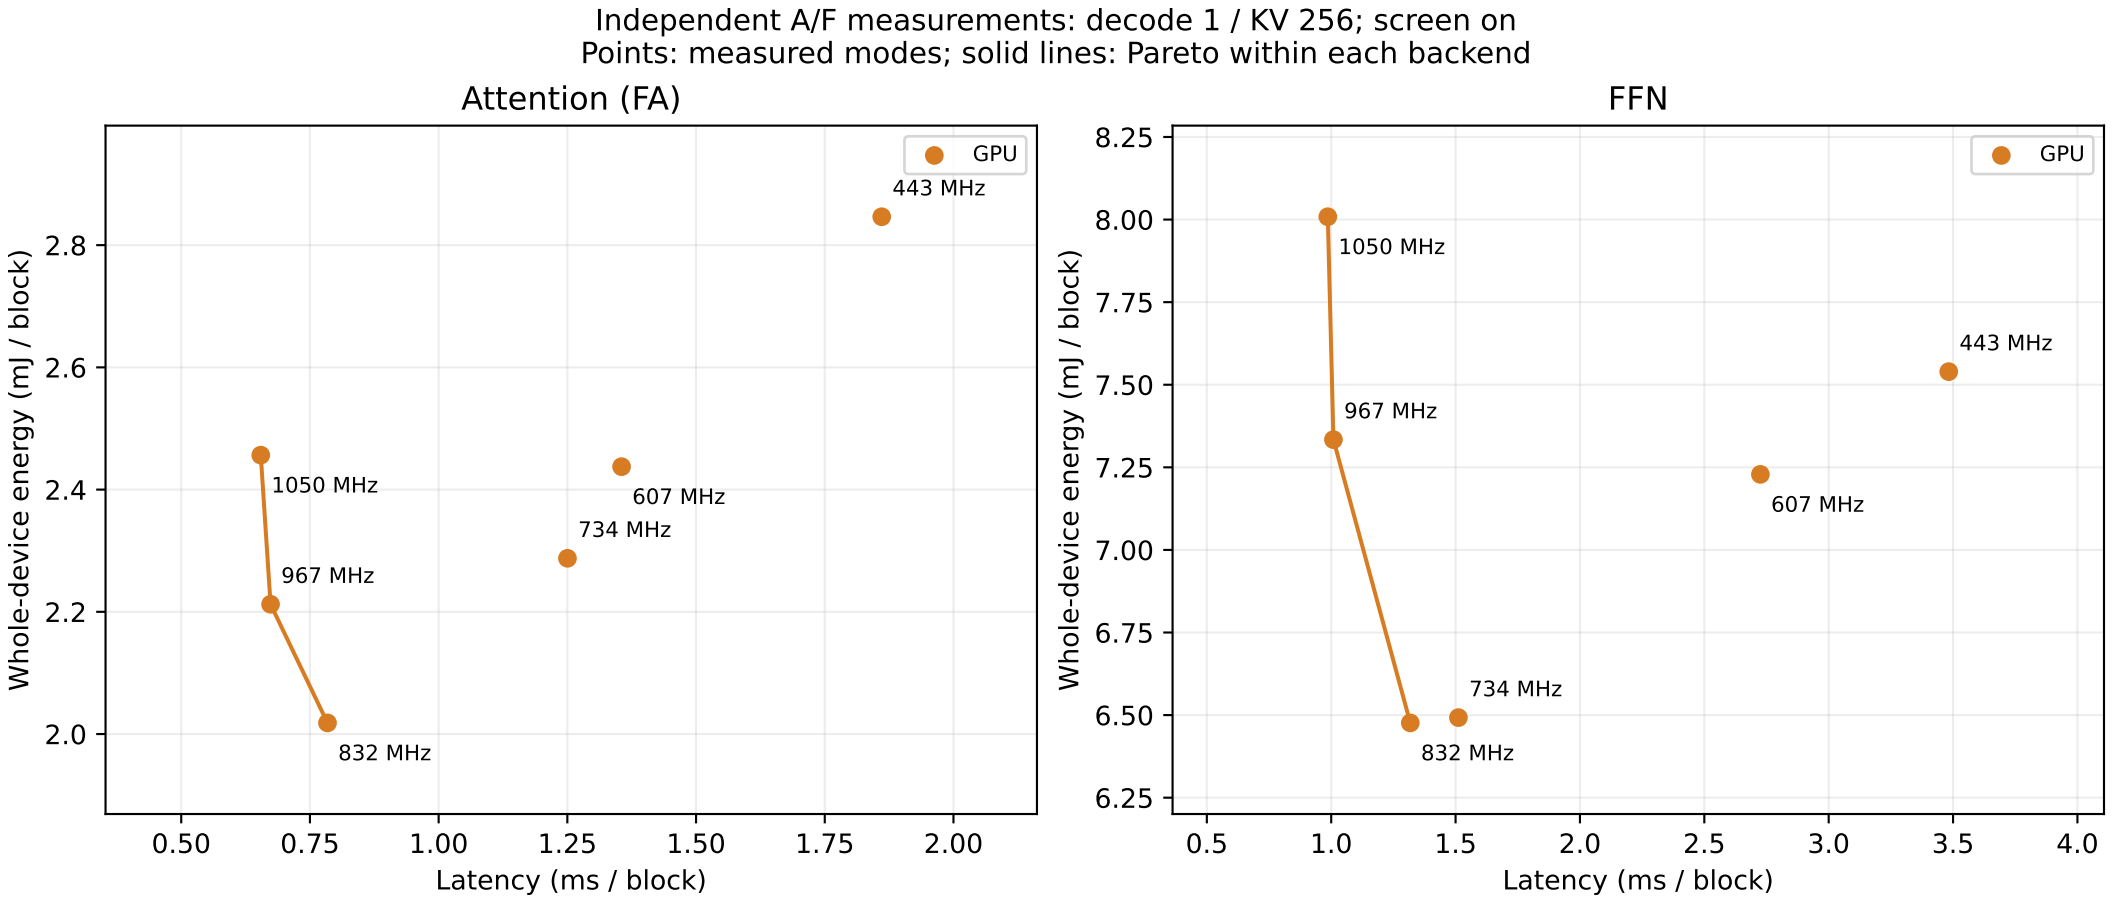
Independent A/F measurements: decode 1 / KV 256; screen on
Points: measured modes; solid lines: Pareto within each backend
Attention (FA)
FFN
GPU
443 MHz
1050 MHz
967 MHz
832 MHz
607 MHz
734 MHz
Whole-device energy (mJ / block)
Latency (ms / block)
2.0
2.2
2.4
2.6
2.8
0.50
0.75
1.00
1.25
1.50
1.75
2.00
6.25
6.50
6.75
7.00
7.25
7.50
7.75
8.00
8.25
0.5
1.0
1.5
2.0
2.5
3.0
3.5
4.0